


Context-dependence (memorisation via attention) at the final layer

Larger Δ = more memorisation encoded in cross-word attention

■ BabyLM (2B tok) ■ Pythia (300B tok)

$\Delta = \text{default } \beta - \text{attn_zeroed } \beta \text{ (mean_pooled, final layer)}$



Corpus-frequency regression: $\text{lm}(|\text{residual}| \sim \log_freq)$ at the final layer, mean_pooled mode.

BabyLM: BabyLM-corpus frequencies; Pythia: Pile frequencies.

Error bars: 95% CI propagated from individual β SEs.